Package 'HEMDAG'

May 21, 2018

Title Hierarchical Ensemble Methods for Directed Acyclic Graphs

Version 2.1.3

Author Marco Notaro [aut, cre] and Giorgio Valentini [aut]
(AnacletoLab, Dipartimento di Informatica, Universita' degli Studi di Milano)

Maintainer Marco Notaro <marco.notaro@unimi.it>

Description An implementation of Hierarchical Ensemble Methods for Directed Acyclic Graphs (DAGs). The 'HEMDAG' package can be used to enhance the predictions of virtually any flat learning methods, by taking into account the hierarchical nature of the classes of a bio-ontology. 'HEMDAG' is specifically designed for exploiting the hierarchical relationships of DAG-structured taxonomies, such as the Human Phenotype Ontology (HPO) or the Gene Ontology (GO), but it can be also safely applied to tree-structured taxonomies (as FunCat), since trees are DAGs. 'HEMDAG' scale nicely both in terms of the complexity of the taxonomy and in the cardinality of the examples. (Marco Notaro, Max Schubach, Peter N. Robinson and Giorgio Valentini (2017) <doi:10.1186/s12859-017-1854-y>).

Depends R (>= 2.10)

License GPL (>= 3)

Encoding UTF-8

Repository CRAN

LazyLoad true

NeedsCompilation no

Imports graph,
RBGL,
PerfMeas,
precrec,
preprocessCore,
methods

Suggests Rgraphviz

RoxygenNote 6.0.1

R topics documented:

HEMDAG-package   *HEMDAG: Hierarchical Ensemble Methods for Directed Acyclic Graphs*

Description

The HEMDAG package provides an implementation of several Hierarchical Ensemble Methods for DAGs. HEMDAG can be used to enhance the predictions of virtually any flat learning methods, by taking into account the hierarchical nature of the classes of a bio-ontology. HEMDAG is specifically designed for exploiting the hierarchical relationships of DAG-structured taxonomies, such as the Human Phenotype Ontology (HPO) or the Gene Ontology (GO), but it can be also safely applied to tree-structured taxonomies (as FunCat), since trees are DAGs. HEMDAG scale nicely both in terms of the complexity of the taxonomy and in the cardinality of the examples.

Details

The HEMDAG package provides many utility functions to handle graph data structures and implements several Hierarchical Ensemble Methods for DAGs:

1. **HTD-DAG**: Hierarchical Top Down (HTD-DAG);

2. **TPR-DAG**: True-Path Rule (TPR-DAG);

3. **DESCENS**: Descendants Ensemble Classifier (DESCENS);

4. **MAX**, **AND**, **OR**: Heuristic Methods, *Obozinski et al.* (Heuristic-Methods);

Author(s)

Marco Notaro and *Giorgio Valentini*, AnacletoLab, DI, Dipartimento di Informatica, Universita' degli Studi di Milano

Maintainer: *Marco Notaro* <marco.notaro@unimi.it>

References

Marco Notaro, Max Schubach, Peter N. Robinson and Giorgio Valentini, *Prediction of Human Phenotype Ontology terms by means of Hierarchical Ensemble methods*, BMC Bioinformatics 2017, 18(1):449, doi:10.1186/s12859-017-1854-y

---

ancestors   *Build ancestors*

Description

Compute the ancestors for each node of a graph

Usage

```
build.ancestors(g)

build.ancestors.per.level(g, levels)

build.ancestors.bottom.up(g, levels)
```

Arguments

g   a graph of class `graphNEL`. It represents the hierarchy of the classes

levels   a list of character vectors. Each component represents a graph level and the elements of any component correspond to nodes. The level 0 coincides with the root node.

Value

`build.ancestos` returns a named list of vectors. Each component corresponds to a node x of the graph and its vector is the set of its ancestors including also x.

`build.ancestors.per.level` returns a named list of vectors. Each component corresponds to a node x of the graph and its vector is the set of its ancestors including also x. The nodes are ordered from root (included) to leaves.

`build.ancestors.bottom.up` a named list of vectors. Each component corresponds to a node x of the graph and its vector is the set of its ancestors including also x. The nodes are ordered from leaves to root (included).

See Also

graph.levels

Examples

```
data(graph);
root <- root.node(g);
anc <- build.ancestors(g);
lev <- graph.levels(g, root=root);
anc.tod <-build.ancestors.per.level(g,lev);
anc.bup <- build.ancestors.bottom.up(g,lev);
```

AUPRC *AUPRC measures*

Description

Function to compute Area under the Precision Recall Curve (AUPRC) through **precrec** package

Usage

```
AUPRC.single.class(target, pred, folds = NULL, seed = NULL)

AUPRC.single.over.classes(labels, predicted, folds = NULL, seed = NULL)
```

Arguments

target	vector of the true labels (0 negative, 1 positive examples)
pred	a numeric vector of the values of the predicted labels (scores)
folds	number of folds on which computing the AUPRC. If folds=NULL (def.), the one-shot AUPRC is computed, otherwise the AUPRC is computed averaged across folds.
seed	initialization seed for the random generator to create folds. Set seed only if folds≠NULL. If seed≠NULL and folds=NULL, the AUPRC averaged across folds is computed without seed initialization.
labels	matrix with the target multilabels: rows correspond to examples and columns to classes. $target[i, j] = 1$ if example i belongs to class j, $target[i, j] = 0$ otherwise.
predicted	a numeric matrix with predicted values (scores): rows correspond to examples and columns to classes.

Details

The AUPRC (for a single class or for a set of classes) is computed either one-shot or averaged across stratified folds.

AUPRC.single.class computes the AUPRC just for a given class.

AUPR.single.over.classes computes the AUPRC for a set of classes, including their average values across all the classes.

For all those classes having zero annotations, the AUPRC is set to 0. These classes are discarded in the computing of the AUPRC averaged across classes, both when the AUPRC is computed one-shot or averaged across stratified folds.

Names of rows and columns of labels and predicted matrix must be provided in the same order, otherwise a stop message is returned

Value

`AUPRC.single.class` returns a numeric value corresponding to the AUPRC for the considered class;

`AUPR.single.over.classes` returns a list with two elements:

1. average: the average AUPRC across classes;

2. per.class: a named vector with AUPRC for each class. Names correspond to classes.

Examples

```
data(labels);
data(scores);
data(graph);
root <- root.node(g);
L <- L[,-which(colnames(L)==root)];
S <- S[,-which(colnames(S)==root)];
PRC.single.class <- AUPRC.single.class(L[,3], S[,3], folds=5, seed=23);
PRC.over.classes <- AUPRC.single.over.classes(L, S, folds=5, seed=23);
```

AUROC   *AUROC measures*

Description

Function to compute the Area under the ROC Curve through **precrec** package

Usage

```
AUROC.single.class(target, pred, folds = NULL, seed = NULL)

AUROC.single.over.classes(labels, predicted, folds = NULL, seed = NULL)
```

Arguments

target	vector of the true labels (0 negative, 1 positive examples)
pred	a numeric vector of the values of the predicted labels (scores)
folds	number of folds on which computing the AUROC. If folds=NULL (def.), the one-shot AUROC is computed, otherwise the AUROC is computed averaged across folds.
seed	initialization seed for the random generator to create folds. Set seed only if folds≠NULL. If seed≠NULL and folds=NULL, the AUROC averaged across folds is computed without seed initialization.
labels	matrix with the target multilabels: rows correspond to examples and columns to classes. $target[i, j] = 1$ if example i belongs to class j, $target[i, j] = 0$ otherwise.
predicted	a numeric matrix with predicted values (scores): rows correspond to examples and columns to classes.

Details

The AUROC (for a single class or for a set of classes) is computed either one-shot or averaged across stratified folds.

`AUROC.single.class` computes the AUROC just for a given class.

`AUROC.single.over.classes` computes the AUROC for a set of classes, including their average values across all the classes.

For all those classes having zero annotations, the AUROC is set to 0.5. These classes are included in the computing of the AUROC averaged across classes, both when the AUROC is computed one-shot or averaged across stratified folds.

The AUROC is set to 0.5 to all those classes having zero annotations. Names of rows and columns of `labels` and `predicted` must be provided in the same order, otherwise a stop message is returned

Value

`AUROC.single.class` returns a numeric value corresponding to the AUROC for the considered class;
`AUPR.single.over.classes` returns a list with two elements:

1. average: the average AUROC across classes;

2. per.class: a named vector with AUROC for each class. Names correspond to classes.

Examples

```
data(labels);
data(scores);
data(graph);
root <- root.node(g);
L <- L[,-which(colnames(L)==root)];
S <- S[,-which(colnames(S)==root)];
AUC.single.class <- AUROC.single.class(L[,3], S[,3], folds=5, seed=23);
AUC.over.classes <- AUROC.single.over.classes(L, S, folds=5, seed=23);
```

---

check.annotation.matrix.integrity
                    Annotation matrix checker

Description

This function assess the integrity of an annotation table in which a transitive closure of annotations was performed

Usage

```
check.annotation.matrix.integrity(anc, ann.spec, ann)
```

Arguments

anc	list of the ancestors of the ontology.
ann.spec	the annotation matrix of the most specific annotations (0/1): rows are genes and columns are terms.
ann	the full annotations matrix (0/1), that is the matrix in which the transitive closure of the annotation was performed. Rows are examples and columns are classes.

Value

If the transitive closure of the annotations is well performed "OK" is returned, otherwise a message error is printed on the stdout

See Also

[build.ancestors](), [transitive.closure.annotations](), [full.annotation.matrix]()

Examples

```
data(graph);
data(labels);
anc <- build.ancestors(g);
tca <- transitive.closure.annotations(L, anc);
check.annotation.matrix.integrity(anc, L, tca);
```

---

check.DAG.integrity   *DAG checker*

Description

This function assess the integrity of a DAG

Usage

```
check.DAG.integrity(g, root = "00")
```

Arguments

g	a graph of class graphNEL. It represents the hierarchy of the classes
root	name of the class that is on the top-level of the hierarchy (def:"00")

Value

If all the nodes are accessible from the root "DAG is OK" is printed, otherwise a message error and the list of the not accessible nodes is printed on the stdout

Examples

```
data(graph);
root <- root.node(g);
check.DAG.integrity(g, root=root);
```

---

children   *Build children*

Description

Compute the children for each node of a graph

Usage

```
build.children(g)

get.children.top.down(g, levels)

get.children.bottom.up(g, levels)
```

Arguments

g   a graph of class `graphNEL`. It represents the hierarchy of the classes

levels   a list of character vectors. Each component represents a graph level and the elements of any component correspond to nodes. The level 0 coincides with the root node.

Value

`build.children` returns a named list of vectors. Each component corresponds to a node x of the graph and its vector is the set of its children

`get.children.top.down` returns a named list of character vectors. Each component corresponds to a node x of the graph (i.e. parent node) and its vector is the set of its children. The nodes are ordered from root (included) to leaves.

`get.children.bottom.up` returns a named list of character vectors. Each component corresponds to a node x of the graph (i.e. parent node) and its vector is the set of its children. The nodes are ordered from leaves (included) to root.

See Also

graph.levels

Examples

```
data(graph);
root <- root.node(g);
children <- build.children(g);
lev <- graph.levels(g, root=root);
children.tod <- get.children.top.down(g,lev);
children.bup <- get.children.bottom.up(g,lev);
```

---

compute.flipped.graph   *Flip Graph*

Description

Compute a directed graph with edges in the opposite direction

Usage

```
compute.flipped.graph(g)
```

Arguments

g   a graphNEL directed graph

Value

a graph (as an object of class graphNEL) with edges in the opposite direction w.r.t. g

Examples

```
data(graph);
g.flipped <- compute.flipped.graph(g);
```

---

constraints.matrix   *Constraints Matrix*

Description

This function returns a matrix with two columns and as many rows as there are edges. The entries of the first columns are the index of the node the edge cames from (i.e. children nodes), the entries of the second columns indicate the index of node the edge is to (i.e. parents nodes). Referring to a DAG this matrix defines a partial order.

Usage

```
constraints.matrix(g)
```

Arguments

g   a graph of class `graphNELL`. It represents the hierarchy of the classes

Value

a constraints matrix w.r.t the graph g

Examples

```
data(graph);
m <- constraints.matrix(g);
```

---

descendants   *Build descendants*

Description

Compute the descendants for each node of a graph

Usage

```
build.descendants(g)

build.descendants.per.level(g, levels)

build.descendants.bottom.up(g, levels)
```

Arguments

g   a graph of class `graphNEL`. It represents the hierarchy of the classes

levels   a list of character vectors. Each component represents a graph level and the elements of any component correspond to nodes. The level 0 coincides with the root node.

Value

`build.descendants` returns a named list of vectors. Each component corresponds to a node x of the graph, and its vector is the set of its descendants including also x.

`build.descendants.per.level` returns a named list of vectors. Each component corresponds to a node x of the graph and its vector is the set of its descendants including also x. The nodes are ordered from root (included) to leaves.

`build.descendants.bottom.up` returns a named list of vectors. Each component corresponds to a node x of the graph and its vector is the set of its descendants including also x. The nodes are ordered from leaves to root (included).

See Also

graph.levels

Examples

```
data(graph);
root <- root.node(g);
desc <- build.descendants(g);
lev <- graph.levels(g, root=root);
desc.tod <- build.descendants.per.level(g,lev);
desc.bup <- build.descendants.bottom.up(g,lev);
```

DESCENS   *DESCENS variants*

Description

The novelty of DESCENS with respect to TPR-DAG algorithm consists in considering the contribution of all the descendants of each node instead of only that of its children, since with the TPR-DAG algorithm the contribution of the descendants of a given node decays exponentially with their distance from the node itself, thus reducing the impact of the predictions made at the most specific levels of the ontology. On the contrary DESCENS predictions are more influenced by the information embedded in the most specific terms of the taxonomy (e.g. leaf nodes), thus putting more emphasis on the terms that most characterize the gene under study.

Usage

```
descens.threshold(S, g, root = "00", t = 0.5)

descens.threshold.free(S, g, root = "00")

descens.weighted.threshold.free(S, g, root = "00", w = 0.5)

descens.weighted.threshold(S, g, root = "00", t = 0.5, w = 0.5)

descens.tau(S, g, root = "00", t = 0.5)
```

Arguments

S	a named flat scores matrix with examples on rows and classes on columns
g	a graph of class graphNEL. It represents the hierarchy of the classes
root	name of the class that it is on the top-level of the hierarchy (def. root="00")
t	threshold for the choice of the positive descendants (def. t=0.5); whereas in the descens.tau variant the parameter t balances the contribution between the positives children of a node i and that of its positives descendants excluding the positives children
w	weight to balance between the contribution of the node i and that of its positive descendants

Details

The *vanilla* DESCENS adopts a per-level bottom-up traversal of the DAG to correct the flat predictions $\hat{y}_i$:

$$\bar{y}_i := \frac{1}{1 + |\Delta_i|}(\hat{y}_i + \sum_{j \in \Delta_i} \bar{y}_j)$$

where Δ_i are the positive descendants of i. Different strategies to select the positive descendants Δ_i can be applied:

1. **Threshold-Free** strategy: as positive descendants we choose those nodes that achieve a score higher than that of their ancestor node i:

$$\Delta_i := \{j \in descendats(i)|\bar{y}_j > \hat{y}_i\}$$

2. **Threshold** strategy: the positive descendants are selected on the basis of a threshold that can ben selected in two different ways:

   (a) for each node a constant threshold $\bar{t}$ is a priori selected:

   $$\phi_i := \{j \in descendats(i)|\bar{y}_j > \bar{t}\}$$

   For instance if the predictions represent probabilities it could be meaningful to a priori select $\bar{t} = 0.5$.

   (b) the threshold is selected to maximize some performance metric $\mathcal{M}$ estimated on the training data, as for instance the F-score or the AUPRC. In other words the threshold is selected to maximize some measure of accuracy of the predictions $\mathcal{M}(j, t)$ on the training data for the class j with respect to the threshold t. The corresponding set of positives $\forall i \in V$ is:

   $$\phi_i := \{j \in descendants(i)|\bar{y}_j > t_j^*, t_j^* = \arg\max_t \mathcal{M}(j, t)\}$$

   For instance t_j^* can be selected from a set of $t \in (0, 1)$ through internal cross-validation techniques.

The weighted DESCENS variants can be simply designed by adding a weight $w \in [0, 1]$ to balance the contribution between the prediction of the classifier associated with the node i and that of its positive descendants:

$$\bar{y}_i := w\hat{y}_i + \frac{(1 - w)}{|\Delta_i|} \sum_{j \in \phi_i} \bar{y}_j$$

The DESCENS-τ variants balances the contribution between the positives children of a node i and that of its positives descendants excluding the children by adding a weight $\tau \in [0, 1]$:

$$\bar{y}_i := \frac{\tau}{1 + |\phi_i|}(\hat{y}_i + \sum_{j \in \phi_i} \bar{y}_j) + \frac{1 - \tau}{1 + |\delta_i|}(\hat{y}_i + \sum_{j \in \delta_i} \bar{y}_j)$$

where ϕ_i are the positive children of i and $\delta_i = \Delta_i \setminus \phi_i$ the descendants of i without its children. If $\tau = 1$ we consider only the contribution of the positive children of i; if $\tau = 0$ only the descendants that are not children contribute to the score, while for intermediate values of τ we can balance the contribution of ϕ_i and δ_i positive nodes.

Value

a named matrix with the scores of the classes corrected according to the DESCENS algorithm.

See Also

[TPR-DAG](#)

Examples

```
data(graph);
data(scores);
data(labels);
root <- root.node(g);
S.descensTF <- descens.threshold.free(S,g,root);
S.descensT <- descens.threshold(S,g,root,t=0.5);
S.descensW <- descens.weighted.threshold.free(S,g,root,w=0.5);
S.descensWT <- descens.weighted.threshold(S,g,root,w=0.5, t=0.5);
S.descensTAU <- descens.tau(S,g,root, t=0.5);
```

---

distances.from.leaves   *Distances from leaves*

Description

This function returns the minimum distance of each node from one of the leaves of the graph

Usage

```
distances.from.leaves(g)
```

Arguments

g   a graph of class graphNEL. It represents the hierarchy of the classes

Value

a named vector. The names are the names of the nodes of the graph g, and their values represent the distance from the leaves. A value equal to 0 is assigned to the leaves, 1 to nodes with distance 1 from a leaf and so on

Examples

```
data(graph);
dist.leaves <- distances.from.leaves(g);
```

---

do.edges.from.HPO.obo *Parse an HPO OBO file*

Description

Read an HPO OBO file (HPO) and write the edges of the DAG on a plain text file.The format of the file is a sequence of rows and each row corresponds to an edge represented through a pair of vertices separated by blanks

Usage

```
do.edges.from.HPO.obo(file = "hp.obo", output.file = "edge.file")
```

Arguments

file   an HPO OBO file

output.file   name of the file of the edges to be written

Value

a text file representing the edges in the format: source destination (i.e. one row for each edge)

Examples

```
## Not run:
hpobo <- "http://purl.obolibrary.org/obo/hp.obo";
do.edges.from.HPO.obo(file=hpobo, output.file="hp.edge");
## End(Not run)
```

Do.flat.scores.normalization
                        Flat scores normalization

Description

High level functions to normalize a flat scores matrix w.r.t. max normalization (MaxNorm) or quantile normalization (Qnorm)

Usage

```
Do.flat.scores.normalization(norm.type = "MaxNorm", flat.file = flat.file,
  flat.dir = flat.dir, flat.norm.dir = flat.norm.dir)
```

Arguments

`norm.type`	can be one of the following two values:

> - MaxNorm (`def.`): each score is divided w.r.t. the max of each class;
> - Qnorm: a quantile normalization is applied. Library preprocessCore is used.

`flat.file`	name of the flat scores matrix (without rda extension)
`flat.dir`	relative path to folder where flat scores matrix is stored
`flat.norm.dir`	the directory where the normalized flat scores matrix must be stored

Details

To apply the quantile normalization the **preprocessCore** library is uded.

Value

the matrix of the scores flat normalized w.r.t. MaxNorm or Qnorm in `flat.norm.dir`

Examples

```
data(scores);
if(!dir.exists("data")){
dir.create("data");
}
if(!dir.exists("results")){
dir.create("results");
}
save(S,file="data/scores.rda");
flat.dir <- "data/";
flat.norm.dir <- "results/";
flat.file <- "scores";
norm.types <- c("MaxNorm","Qnorm");
for(norm.type in norm.types){
Do.flat.scores.normalization(norm.type=norm.type, flat.file=flat.file,
flat.dir=flat.dir, flat.norm.dir=flat.norm.dir);
}
```

`Do.full.annotation.matrix`

Do full annotations matrix

Description

High-level function to obtain a full annotation matrix, that is a matrix in which the transitive closure of annotations was performed, respect to a given weighted adiacency matrix

Usage

```
Do.full.annotation.matrix(anc.file.name = anc.file.name, anc.dir = anc.dir,
  net.file = net.file, net.dir = net.dir, ann.file.name = ann.file.name,
  ann.dir = ann.dir, output.name = output.name, output.dir = output.dir)
```

Arguments

`anc.file.name`	name of the file containing the list for each node the list of all its ancestor (without `rda` extension)
`anc.dir`	relative path to directory where the ancestor file is stored
`net.file`	name of the file containing the weighted adjiacency matrix of the graph (without `rda` extension)
`net.dir`	relative path to directory where the weighted adjiacency matrix is stored
`ann.file.name`	name of the file containing the matrix of the most specific annotations (without `rda` extension)
`ann.dir`	relative path to directory where the matrix of the most specific annotation is stored
`output.name`	name of the output file without rda extension (without `rda` extension)
`output.dir`	relative path to directory where the output file must be stored

Value

a full annotation matrix T, that is a matrix in which the transitive closure of annotations was performed. Rows correspond to genes of the input weighted adjiacency matrix and columns to terms. $T[i,j] = 1$ means that gene i is annotated for the term j, $T[i,j] = 0$ means that gene i is not annotated for the term j.

See Also

`full.annotation.matrix`

Examples

```
data(graph);
data(labels);
data(wadj);
if (!dir.exists("data")){
dir.create("data");
}
if (!dir.exists("results")){
dir.create("results");
}
anc <- build.ancestors(g);
save(anc,file="data/ancestors.rda");
save(g,file="data/graph.rda");
save(L,file="data/labels.rda");
save(W,file="data/wadj.rda");
anc.dir <- net.dir <- ann.dir <- "data/";
```

```
output.dir <- "results/";
anc.file.name <- "ancestors";
net.file <- "wadj";
ann.file.name <- "labels";
output.name <- "full.ann.matrix";
Do.full.annotation.matrix(anc.file.name=anc.file.name, anc.dir=anc.dir, net.file=net.file,
net.dir=net.dir, ann.file.name=ann.file.name, ann.dir=ann.dir, output.name=output.name,
output.dir=output.dir);
```

Do.heuristic.methods   *Do Heuristic Methods*

Description

High level function to compute the hierarchical heuristic methods MAX, AND, OR (Heuristic
Methods MAX, AND, OR (*Obozinski et al., Genome Biology, 2008*)

Usage

```
Do.heuristic.methods(heuristic.fun = "AND", norm = TRUE, norm.type = NULL,
  folds = 5, seed = 23, n.round = 3, f.criterion = "F",
  rec.levels = seq(from = 0.1, to = 1, by = 0.1), flat.file = flat.file,
  ann.file = ann.file, dag.file = dag.file, flat.dir = flat.dir,
  ann.dir = ann.dir, dag.dir = dag.dir, hierScore.dir = hierScore.dir,
  perf.dir = perf.dir)
```

Arguments

heuristic.fun   can be one of the following three values:

1. "MAX": run the heuristic method MAX;
2. "AND": run the heuristic method AND;
3. "OR": run the heuristic method OR;

norm   boolean value:

- TRUE (def.): the flat scores matrix has been already normalized in according
  to a normalization method;
- FALSE: the flat scores matrix has not been normalized yet. See the parameter
  norm.type for which normalization can be applied.

norm.type   can be one of the following three values:

1. NULL (def.): set norm.type to NULL if and only if the parameter norm is set
   to TRUE;
2. MaxNorm: each score is divided for the maximum of each class;
3. Qnorm: quantile normalization. **preprocessCore** package is used.

folds   number of folds of the cross validation on which computing the performance
                 metrics averaged across folds (def. 5). If folds=NULL, the performance metrics
                 are computed one-shot, otherwise the performance metrics are averaged across
                 folds.

seed	initialization seed for the random generator to create folds (def. 23). If NULL folds are generated without seed initialization.
n.round	number of rounding digits to be applied to the hierarchical scores matrix (def. 3). It is used for choosing the best threshold on the basis of the best F-measure
f.criterion	character. Type of F-measure to be used to select the best F-measure. Two possibilities:

        1. F (def.): corresponds to the harmonic mean between the average precision and recall

        2. avF: corresponds to the per-example F-score averaged across all the examples

rec.levels	a vector with the desired recall levels (def: from:0.1, to:0.9, by:0.1) to compute the the Precision at fixed Recall level (PXR)
flat.file	name of the file containing the flat scores matrix to be normalized or already normalized (without rda extension)
ann.file	name of the file containing the the label matrix of the examples (without rda extension)
dag.file	name of the file containing the graph that represents the hierarchy of the classes (without rda extension)
flat.dir	relative path where flat scores matrix is stored
ann.dir	relative path where annotation matrix is stored
dag.dir	relative path where graph is stored
hierScore.dir	relative path where the hierarchical scores matrix must be stored
perf.dir	relative path where all the performance measures must be stored

Details

The function checks if the number of classes between the flat scores matrix and the annotations matrix mismatched. If so, the number of terms of the annotations matrix is shrunk to the number of terms of the flat scores matrix and the corresponding subgraph is computed as well. N.B.: it is supposed that all the nodes of the subgraph are accessible from the root.

Value

Two rda files stored in the respective output directories:

1. Hierarchical Scores Results: a matrix with examples on rows and classes on columns representing the computed hierarchical scores for each example and for each considered class. It is stored in the hierScore.dir directory.

2. Performance Measures: *flat* and *hierarchical* performace results:

    (a) AUPRC results computed though AUPRC.single.over.classes (AUPRC);

    (b) AUROC results computed through AUROC.single.over.classes (AUROC);

    (c) PXR results computed though PXR.at.multiple.recall.levels.over.classes (PXR);

    (d) FMM results computed though compute.Fmeasure.multilabel (FMM);

It is stored in the perf.dir directory.

See Also

Heuristic-Methods

Examples

```
data(graph);
data(scores);
data(labels);
if (!dir.exists("data")){
dir.create("data");
}
if (!dir.exists("results")){
dir.create("results");
}
save(g,file="data/graph.rda");
save(L,file="data/labels.rda");
save(S,file="data/scores.rda");
dag.dir <- flat.dir <- ann.dir <- "data/";
hierScore.dir <- perf.dir <- "results/";
rec.levels <- seq(from=0.1, to=1, by=0.1);
dag.file <- "graph";
flat.file <- "scores";
ann.file <- "labels";
Do.heuristic.methods(heuristic.fun="AND", norm=FALSE, norm.type="MaxNorm",
folds=5, seed=23, n.round=3, f.criterion ="F", rec.levels=rec.levels,
flat.file=flat.file, ann.file=ann.file, dag.file=dag.file, flat.dir=flat.dir,
ann.dir=ann.dir, dag.dir=dag.dir, hierScore.dir=hierScore.dir, perf.dir=perf.dir);
```

```
Do.heuristic.methods.holdout
```
                    Do Heuristic Methods holdout

Description

High level function to compute the hierarchical heuristic methods MAX, AND, OR (Heuristic Methods MAX, AND, OR (*Obozinski et al., Genome Biology, 2008*) applying a classical hold-out procedure

Usage

```
Do.heuristic.methods.holdout(heuristic.fun = "AND", norm = TRUE,
  norm.type = NULL, folds = 5, seed = 23, n.round = 3,
  f.criterion = "F", rec.levels = seq(from = 0.1, to = 1, by = 0.1),
  flat.file = flat.file, ann.file = ann.file, dag.file = dag.file,
  ind.test.set = ind.test.set, ind.dir = ind.dir, flat.dir = flat.dir,
  ann.dir = ann.dir, dag.dir = dag.dir, hierScore.dir = hierScore.dir,
  perf.dir = perf.dir)
```

Arguments

heuristic.fun	can be one of the following three values:

          1. "MAX": run the heuristic method MAX;

          2. "AND": run the heuristic method AND;

          3. "OR": run the heuristic method OR;

norm   boolean value:

- TRUE (def.): the flat scores matrix has been already normalized in according to a normalization method;
- FALSE: the flat scores matrix has not been normalized yet. See the parameter norm for which normalization can be applied.

norm.type   can be one of the following three values:

          1. NONE (def.): set norm.type to NONE if and only if the parameter norm is set to TRUE;

          2. MaxNorm: each score is divided for the maximum of each class;

          3. Qnorm: quantile normalization. **preprocessCore** package is used.

folds   number of folds of the cross validation on which computing the performance metrics averaged across folds (def. 5). If folds=NULL, the performance metrics are computed one-shot, otherwise the performance metrics are averaged across folds.

seed   initialization seed for the random generator to create folds (def. 23). If NULL folds are generated without seed initialization.

n.round   number of rounding digits to be applied to the hierarchical scores matrix (def. 3). It is used for choosing the best threshold on the basis of the best F-measure

f.criterion   character. Type of F-measure to be used to select the best F-measure. Two possibilities:

          1. F (def.): corresponds to the harmonic mean between the average precision and recall

          2. avF: corresponds to the per-example F-score averaged across all the examples

rec.levels   a vector with the desired recall levels (def: from:0.1, to:0.9, by:0.1) to compute the the Precision at fixed Recall level (PXR)

flat.file   name of the file containing the flat scores matrix to be normalized or already normalized (without rda extension)

ann.file   name of the file containing the the label matrix of the examples (without rda extension)

dag.file   name of the file containing the graph that represents the hierarchy of the classes (without rda extension)

ind.test.set   name of the file containing a vector of integer numbers corresponding to the indices of the elements (rows) of scores matrix to be used in the test set

ind.dir   relative path to folder where ind.test.set is stored

flat.dir   relative path where flat scores matrix is stored

ann.dir   relative path where annotation matrix is stored

`dag.dir`	relative path where graph is stored
`hierScore.dir`	relative path where the hierarchical scores matrix must be stored
`perf.dir`	relative path where the term-centric and protein-centric measures must be stored

Details

The function checks if the number of classes between the flat scores matrix and the annotations matrix mismatched. If so, the number of terms of the annotations matrix is shrunk to the number of terms of the flat scores matrix and the corresponding subgraph is computed as well. N.B.: it is supposed that all the nodes of the subgraph are accessible from the root.

Value

Two rda files stored in the respective output directories:

1. `Hierarchical Scores Results`: a matrix with examples on rows and classes on columns representing the computed hierarchical scores for each example and for each considered class. It is stored in the `hierScore.dir` directory.

2. `Performance Measures`: *flat* and *hierarchical* performace results:

   (a) AUPRC results computed though AUPRC.single.over.classes ([AUPRC](#));
   (b) AUROC results computed through AUROC.single.over.classes ([AUROC](#));
   (c) PXR results computed though PXR.at.multiple.recall.levels.over.classes ([PXR](#));
   (d) FMM results computed though compute.Fmeasure.multilabel ([FMM](#));

It is stored in the `perf.dir` directory.

Examples

```
data(graph);
data(scores);
data(labels);
data(test.index);
if (!dir.exists("data")){
dir.create("data");
}
if (!dir.exists("results")){
dir.create("results");
}
save(g,file="data/graph.rda");
save(L,file="data/labels.rda");
save(S,file="data/scores.rda");
save(test.index, file="data/test.index.rda");
ind.dir <- dag.dir <- flat.dir <- ann.dir <- "data/";
hierScore.dir <- perf.dir <- "results/";
rec.levels <- seq(from=0.1, to=1, by=0.1);
ind.test.set <- "test.index";
dag.file <- "graph";
flat.file <- "scores";
ann.file <- "labels";
Do.heuristic.methods.holdout(heuristic.fun="MAX", norm=FALSE, norm.type="MaxNorm",
```

```
folds=NULL, seed=23, n.round=3, f.criterion ="F", rec.levels=rec.levels,
flat.file=flat.file, ann.file=ann.file, dag.file=dag.file,
ind.test.set=ind.test.set, ind.dir=ind.dir, flat.dir=flat.dir, ann.dir=ann.dir,
dag.dir=dag.dir, hierScore.dir=hierScore.dir, perf.dir=perf.dir);
```

Do.HTD   *HTD-DAG vanilla*

Description

High level function to correct the computed scores in a hierarchy according to the HTD-DAG algorithm

Usage

```
Do.HTD(norm = TRUE, norm.type = NULL, folds = 5, seed = 23,
  n.round = 3, f.criterion = "F", rec.levels = seq(from = 0.1, to = 1, by
  = 0.1), flat.file = flat.file, ann.file = ann.file, dag.file = dag.file,
  flat.dir = flat.dir, ann.dir = ann.dir, dag.dir = dag.dir,
  hierScore.dir = hierScore.dir, perf.dir = perf.dir)
```

Arguments

norm
: boolean value:

  - TRUE (def.): the flat scores matrix has been already normalized in according to a normalization method;
  - FALSE: the flat scores matrix has not been normalized yet. See the parameter norm.type for which normalization can be applied.

norm.type
: can be one of the following three values:

  1. NULL (def.): set norm.type to NULL if and only if the parameter norm is set to TRUE;
  2. MaxNorm: each score is divided for the maximum of each class;
  3. Qnorm: quantile normalization. **preprocessCore** package is used.

folds
: number of folds of the cross validation on which computing the performance metrics averaged across folds (def. 5). If folds=NULL, the performance metrics are computed one-shot, otherwise the performance metrics are averaged across folds.

seed
: initialization seed for the random generator to create folds (def. 23). If NULL folds are generated without seed initialization.

n.round
: number of rounding digits to be applied to the hierarchical scores matrix (def. 3). It is used for choosing the best threshold on the basis of the best F-measure

f.criterion
: character. Type of F-measure to be used to select the best F-measure. Two possibilities:

  1. F (def.): corresponds to the harmonic mean between the average precision and recall

      2. `avF`: corresponds to the per-example `F-score` averaged across all the examples

`rec.levels`	a vector with the desired recall levels (def: `from:0.1, to:0.9, by:0.1`) to compute the the Precision at fixed Recall level (PXR)
`flat.file`	name of the file containing the flat scores matrix to be normalized or already normalized (without rda extension)
`ann.file`	name of the file containing the the label matrix of the examples (without rda extension)
`dag.file`	name of the file containing the graph that represents the hierarchy of the classes (without rda extension)
`flat.dir`	relative path where flat scores matrix is stored
`ann.dir`	relative path where annotation matrix is stored
`dag.dir`	relative path where graph is stored
`hierScore.dir`	relative path where the hierarchical scores matrix must be stored
`perf.dir`	relative path where all the performance measures must be stored

Details

The function checks if the number of classes between the flat scores matrix and the annotations matrix mismatched. If so, the number of terms of the annotations matrix is shrunk to the number of terms of the flat scores matrix and the corresponding subgraph is computed as well. N.B.: it is supposed that all the nodes of the subgraph are accessible from the root.

Value

Two `rda` files stored in the respective output directories:

1. `Hierarchical Scores Results`: a matrix with examples on rows and classes on columns representing the computed hierarchical scores for each example and for each considered class. It is stored in the `hierScore.dir` directory.

2. `Performance Measures`: *flat* and *hierarchical* performace results:

    (a) AUPRC results computed though `AUPRC.single.over.classes` ([AUPRC](#));

    (b) AUROC results computed through `AUROC.single.over.classes` ([AUROC](#));

    (c) PXR results computed though `PXR.at.multiple.recall.levels.over.classes` ([PXR](#));

    (d) FMM results computed though `compute.Fmeasure.multilabel` ([FMM](#));

It is stored in the `perf.dir` directory.

See Also

[HTD-DAG](#)

Examples

```
data(graph);
data(scores);
data(labels);
if (!dir.exists("data")){
dir.create("data");
}
if (!dir.exists("results")){
dir.create("results");
}
save(g,file="data/graph.rda");
save(L,file="data/labels.rda");
save(S,file="data/scores.rda");
dag.dir <- flat.dir <- ann.dir <- "data/";
hierScore.dir <- perf.dir <- "results/";
rec.levels <- seq(from=0.1, to=1, by=0.1);
dag.file <- "graph";
flat.file <- "scores";
ann.file <- "labels";
Do.HTD(norm=FALSE, norm.type="MaxNorm", folds=5, seed=23, n.round=3,
f.criterion="F", rec.levels=rec.levels, flat.file=flat.file, ann.file=ann.file,
dag.file=dag.file, flat.dir=flat.dir, ann.dir=ann.dir, dag.dir=dag.dir,
hierScore.dir=hierScore.dir, perf.dir=perf.dir);
```

Do.HTD.holdout   *HTD-DAG holdout*

Description

High level function to correct the computed scores in a hierarchy according to the HTD-DAG algorithm applying a classical holdout procedure

Usage

```
Do.HTD.holdout(norm = TRUE, norm.type = NULL, folds = 5, seed = 23,
  n.round = 3, f.criterion = "F", rec.levels = seq(from = 0.1, to = 1, by
  = 0.1), flat.file = flat.file, ann.file = ann.file, dag.file = dag.file,
  ind.test.set = ind.test.set, ind.dir = ind.dir, flat.dir = flat.dir,
  ann.dir = ann.dir, dag.dir = dag.dir, hierScore.dir = hierScore.dir,
  perf.dir = perf.dir)
```

Arguments

norm   boolean value:

- TRUE (def.): the flat scores matrix has been already normalized in according to a normalization method;
- FALSE: the flat scores matrix has not been normalized yet. See the parameter norm for which normalization can be applied.

norm.type	can be one of the following three values:

1. NULL (def.): set norm.type to NULL if and only if the parameter norm is set to TRUE;

2. MaxNorm: each score is divided for the maximum of each class;

3. Qnorm: quantile normalization. **preprocessCore** package is used.

folds	number of folds of the cross validation on which computing the performance metrics averaged across folds (def. 5). If folds=NULL, the performance metrics are computed one-shot, otherwise the performance metrics are averaged across folds.
seed	initialization seed for the random generator to create folds (def. 23). If NULL folds are generated without seed initialization.
n.round	number of rounding digits to be applied to the hierarchical scores matrix (def. 3). It is used for choosing the best threshold on the basis of the best F-measure
f.criterion	character. Type of F-measure to be used to select the best F-measure. Two possibilities:

1. F (def.): corresponds to the harmonic mean between the average precision and recall

2. avF: corresponds to the per-example F-score averaged across all the examples

rec.levels	a vector with the desired recall levels (def: from:0.1, to:0.9, by:0.1) to compute the the Precision at fixed Recall level (PXR)
flat.file	name of the file containing the flat scores matrix to be normalized or already normalized (without rda extension)
ann.file	name of the file containing the the label matrix of the examples (without rda extension)
dag.file	name of the file containing the graph that represents the hierarchy of the classes (without rda extension)
ind.test.set	name of the file containing a vector of integer numbers corresponding to the indices of the elements (rows) of scores matrix to be used in the test set
ind.dir	relative path to folder where ind.test.set is stored
flat.dir	relative path where flat scores matrix is stored
ann.dir	relative path where annotation matrix is stored
dag.dir	relative path where graph is stored
hierScore.dir	relative path where the hierarchical scores matrix must be stored
perf.dir	relative path where all the performance measures must be stored

Details

The function checks if the number of classes between the flat scores matrix and the annotations matrix mismatched. If so, the number of terms of the annotations matrix is shrunk to the number of terms of the flat scores matrix and the corresponding subgraph is computed as well. N.B.: it is supposed that all the nodes of the subgraph are accessible from the root.

Value

Two `rda` files stored in the respective output directories:

1. `Hierarchical Scores Results`: a matrix with examples on rows and classes on columns representing the computed hierarchical scores for each example and for each considered class. It is stored in the `hierScore.dir` directory.

2. `Performance Measures`: *flat* and *hierarchical* performace results:

   (a) AUPRC results computed though `AUPRC.single.over.classes` ([AUPRC](#));
   (b) AUROC results computed through `AUROC.single.over.classes` ([AUROC](#));
   (c) PXR results computed though `PXR.at.multiple.recall.levels.over.classes` ([PXR](#));
   (d) FMM results computed though `compute.Fmeasure.multilabel` ([FMM](#));

It is stored in the `perf.dir` directory.

See Also

[HTD-DAG](#)

Examples

```
data(graph);
data(scores);
data(labels);
data(test.index);
if (!dir.exists("data")){
dir.create("data");
}
if (!dir.exists("results")){
dir.create("results");
}
save(g,file="data/graph.rda");
save(L,file="data/labels.rda");
save(S,file="data/scores.rda");
save(test.index, file="data/test.index.rda");
ind.dir <- dag.dir <- flat.dir <- ann.dir <- "data/";
hierScore.dir <- perf.dir <- "results/";
rec.levels <- seq(from=0.1, to=1, by=0.1);
ind.test.set <- "test.index";
dag.file <- "graph";
flat.file <- "scores";
ann.file <- "labels";
Do.HTD.holdout(norm=FALSE, norm.type="MaxNorm", n.round=3, f.criterion ="F", folds=NULL, seed=23,
rec.levels=rec.levels, flat.file=flat.file, ann.file=ann.file, dag.file=dag.file,
ind.test.set=ind.test.set, ind.dir=ind.dir, flat.dir=flat.dir, ann.dir=ann.dir, dag.dir=dag.dir,
hierScore.dir=hierScore.dir, perf.dir=perf.dir);
```

---

do.subgraph   *Build subgraph*

Description

This function returns a subgraph with only the supplied nodes and any edges between them

Usage

```
do.subgraph(nd, g, edgemode = "directed")
```

Arguments

nd   a vector with the nodes for which the subgraph must be built

g   a graph of class graphNEL. It represents the hierarchy of the classes

edgemode   can be "directed" or "undirected"

Value

a subgraph with only the supplied nodes

Examples

```
data(graph);
anc <- build.ancestors(g);
nd <- anc[["HP:0001371"]];
subg <- do.subgraph(nd, g, edgemode="directed");
```

---

do.submatrix   *Build submatrix*

Description

Terms having less than n annotations are pruned. Terms having exactly n annotations are discarded
as well.

Usage

```
do.submatrix(ann, n)
```

Arguments

ann   the annotations matrix (0/1). Rows are examples and columns are classes

n   integer number of annotations to be pruned

Value

Matrix of annotations having only those terms with more than n annotations

Examples

```
data(labels);
subm <- do.submatrix(L,5);
```

```
do.unstratified.cv.data
```
Unstratified cross-validation

Description

This function splits a dataset in k-fold in an unstratified way (that is a fold may not have an equal amount of positive and negative examples). This function is used to perform k-fold cross-validation experiments in a hierarchical correction contest where splitting dataset in a stratified way is not needed.

Usage

```
do.unstratified.cv.data(S, kk = 5, seed = NULL)
```

Arguments

S   matrix of the flat scores. It must be a named matrix, where rows are example (e.g. genes) and columns are classes/terms (e.g. HPO terms)

kk   number of folds in which to split the dataset (def. k=5)

seed   seed for the random generator. If NULL (def.) no initialization is performed

Value

a list with $k = kk$ components (folds). Each component of the list is a character vector contains the index of the examples, i.e. the index of the rows of the matrix S

Examples

```
data(scores);
foldIndex <- do.unstratified.cv.data(S, kk=5, seed=23);
```

---

example.datasets *Small real example datasets*

Description

Collection of real sub-datasets used in the examples of the **HEMDAG** package

Usage

```
data(graph)
data(labels)
data(scores)
data(wadj)
data(test.index)
```

Details

The DAG g contained in graph data is an object of class graphNEL. The graph g has 23 nodes and 30 edges and represents the "ancestors view" of the HPO term *Camptodactyly of finger* ("HP:0100490").

The matrix L contained in the labels data is a 100 X 23 matrix, whose rows correspondes to genes (*Entrez GeneID*) and columns to HPO classes. $L[i, j] = 1$ means that the gene i belong to class j, $L[i, j] = 0$ means that the gene i does not belong to class j. The classes of the matrix L correspond to the nodes of the graph g.

The matrix S contained in the scores data is a named 100 X 23 flat scores matrix, representing the likelihood that a given gene belongs to a given class: higher the value higher the likelihood. The classes of the matrix S correspond to the nodes of the graph g.

The matrix W contained in the wadj data is a named 100 X 100 symmetric weighted adjacency matrix, whose rows and columns correspond to genes.The genes names (*Entrez GeneID*) of the adjacency matrix W correspond to the genes names of the flat scores matrix S and to genes names of the target multilabel matrix L.

The vector of integer numbers test.index contained in the test.index data refers to the index of the examples of the scores matrix S to be used in the test set. It is useful only in holdout experiments.

Note

Some examples of full data sets for the prediction of HPO terms are available at the following link. Note that the processing of the full datasets should be done similarly to the processing of the small data examples provided directly in this package. Please read the README clicking the link above to know more details about the available full datasets.

---

find.best.f   *Best hierarchical F-score*

Description

Function to select the best hierarchical F-score by choosing an appropriate threshold in the scores

Usage

```
find.best.f(target, pred, n.round = 3, f.criterion = "F", verbose = TRUE,
  b.per.example = FALSE)
```

Arguments

target
: matrix with the target multilabels: rows correspond to examples and columns to classes. $target[i, j] = 1$ if example i belongs to class j, $target[i, j] = 0$ otherwise

pred
: a numeric matrix with continuous predicted values (scores): rows correspond to examples and columns to classes

n.round
: number of rounding digits to be applied to pred (default=3)

f.criterion
: character. Type of F-measure to be used to select the best F-score. There are two possibilities:

  1. F (def.) corresponds to the harmonic mean between the average precision and recall;
  2. avF corresponds to the per-example F-score averaged across all the examples.

verbose
: boolean. If TRUE (def.) the number of iterations are printed on stdout

b.per.example
: boolean.

  - TRUE: results are returned for each example;
  - FALSE: only the average results are returned

Details

All the examples having no positive annotations are discarded. The predicted scores matrix (pred) is rounded according to parameter n.round and all the values of pred are divided by max(pred). Then all the thresholds corresponding to all the different values included in pred are attempted, and the threshold leading to the maximum F-measure is selected.

Names of rows and columns of target and pred matrix must be provided in the same order, otherwise a stop message is returned

Value

Two different outputs respect to the input parameter `b.per.example`:

- `b.per.example==FALSE`: a list with a single element average. A named vector with 7 elements relative to the best result in terms of the F.measure: Precision (P), Recall (R), Specificity (S), F.measure (F), av.F.measure (av.F), Accuracy (A) and the best selected Threshold (T). F is the F-measure computed as the harmonic mean between the average precision and recall; av.F is the F-measure computed as the average across examples and T is the best selected threshold;

- `b.per.example==FALSE`: a list with two elements:

  1. average: a named vector with with 7 elements relative to the best result in terms of the F.measure: Precision (P), Recall (R), Specificity (S), F.measure (F), av.F.measure (av.F), Accuracy (A) and the best selected Threshold (T).
  2. per.example: a named matrix with the Precision (P), Recall (R), Specificity (S), Accuracy (A), F-measure (F), av.F-measure (av.F) and the best selected Threshold (T) for each example. Row names correspond to examples, column names correspond respectively to Precision (P), Recall (R), Specificity (S), Accuracy (A), F-measure (F), av.F-measure (av.F) and the best selected Threshold (T).

Examples

```
data(graph);
data(labels);
data(scores);
root <- root.node(g);
L <- L[,-which(colnames(L)==root)];
S <- S[,-which(colnames(S)==root)];
FMM <- find.best.f(L, S, n.round=3, f.criterion="F", verbose=TRUE, b.per.example=TRUE);
```

---

find.leaves   *Leaves*

Description

Find the leaves of a directed graph

Usage

```
find.leaves(g)
```

Arguments

g   a graph of class `graphNEL`. It represents the hierarchy of the classes

Value

a vector with the names of the leaves of g

Examples

```
data(graph);
leaves <- find.leaves(g);
```

FMM   *Compute Multilabel F-measure*

Description

Function to compute the best hierarchical F-score either one-shot or averaged across folds

Usage

```
compute.Fmeasure.multilabel(target, pred, n.round = 3, f.criterion = "F",
  verbose = TRUE, b.per.example = FALSE, folds = NULL, seed = NULL)
```

Arguments

target	matrix with the target multilabels: rows correspond to examples and columns to classes. $target[i, j] = 1$ if example i belongs to class j, $target[i, j] = 0$ otherwise
pred	a numeric matrix with predicted values (scores): rows correspond to examples and columns to classes
n.round	number of rounding digits to be applied to pred (default=3)
f.criterion	character. Type of F-measure to be used to select the best F-score. There are two possibilities:

     1. F (def.) corresponds to the harmonic mean between the average precision and recall;

     2. avF corresponds to the per-example F-score averaged across all the examples.

verbose	boolean. If TRUE (def.) the number of iterations are printed on stdout
b.per.example	boolean.

- TRUE: results are returned for each example;
- FALSE: only the average results are returned

folds	number of folds on which computing the AUROC. If folds=NULL (def.), the one-shot AUROC is computed, otherwise the AUROC is computed averaged across folds.
seed	initialization seed for the random generator to create folds. Set seed only if folds≠NULL. If seed≠NULL and folds=NULL, the AUROC averaged across folds is computed without seed initialization.

Details

Names of rows and columns of target and pred matrix must be provided in the same order, otherwise a stop message is returned

Value

Two different outputs respect to the input parameter `b.per.example`:

- `b.per.example==FALSE`: a list with a single element average. A named vector with 7 elements relative to the best result in terms of the F.measure: Precision (P), Recall (R), Specificity (S), F.measure (F), av.F.measure (av.F), Accuracy (A) and the best selected Threshold (T). F is the F-measure computed as the harmonic mean between the average precision and recall; av.F is the F-measure computed as the average across examples and T is the best selected threshold;

- `b.per.example==FALSE`: a list with two elements:

  1. average: a named vector with with 7 elements relative to the best result in terms of the F.measure: Precision (P), Recall (R), Specificity (S), F.measure (F), av.F.measure (av.F), Accuracy (A) and the best selected Threshold (T).

  2. per.example: a named matrix with the Precision (P), Recall (R), Specificity (S), Accuracy (A), F-measure (F), av.F-measure (av.F) and the best selected Threshold (T) for each example. Row names correspond to examples, column names correspond respectively to Precision (P), Recall (R), Specificity (S), Accuracy (A), F-measure (F), av.F-measure (av.F) and the best selected Threshold (T).

Examples

```
data(graph);
data(labels);
data(scores);
root <- root.node(g);
L <- L[,-which(colnames(L)==root)];
S <- S[,-which(colnames(S)==root)];
FMM <- compute.Fmeasure.multilabel(L, S, n.round=3, f.criterion="F", verbose=TRUE,
b.per.example=TRUE, folds=5, seed=23);
```

`full.annotation.matrix`

Full annotations matrix

Description

Construct a full annotations table using ancestors and the most specific annotations table w.r.t. a given weighted adjacency matrix (wadj). The rows of the full annotations matrix correspond to all the examples of the given weighted adjacency matrix and the columns to the class/terms. The transitive closure of the annotations is performed.

Usage

```
full.annotation.matrix(W, anc, ann.spec)
```

Arguments

W	symmetric adjacency weighted matrix of the graph
anc	list of the ancestors of the ontology.
ann.spec	the annotation matrix of the most specific annotations (0/1): rows are genes and columns are classes.

Details

The examples present in the annotation matrix (`ann.spec`) but not in the adjacency weighted matrix (`W`) are purged.

Value

a full annotation table T, that is a matrix in which the transitive closure of annotations was performed. Rows correspond to genes of the weighted adjiacency matrix and columns to terms. $T[i, j] = 1$ means that gene i is annotated for the term j, $T[i, j] = 0$ means that gene i is not annotated for the term j.

See Also

weighted.adjacency.matrix, build.ancestors,
specific.annotation.matrix, transitive.closure.annotations

Examples

```
data(wadj);
data(graph);
data(labels);
anc <- build.ancestors(g);
full.ann <- full.annotation.matrix(W, anc, L);
```

---

graph.levels *Build Graph Levels*

Description

This function groups a set of nodes in according to their maximum depth in the graph. It first inverts the weights of the graph and then applies the Bellman Ford algorithm to find the shortest path, achieving in this way the longest path

Usage

```
graph.levels(g, root = "00")
```

Arguments

g	an object of class graphNEL
root	name of the root node (def. root="00")

Value

a list of the nodes grouped w.r.t. the distance from the root: the first element of the list corresponds to the root node (level 0), the second to nodes at maximum distance 1 (level 1), the third to the node at maximum distance 3 (level 2) and so on.

Examples

```
data(graph);
root <- root.node(g);
lev <- graph.levels(g, root=root);
```

Heuristic-Methods   *Obozinski Heuristic Methods*

Description

Implementation of the Heuristic Methods MAX, AND, OR (*Obozinski et al., Genome Biology, 2008, doi:10.1186/gb-2008-9-s1-s6*)

Usage

```
heuristicMAX(S, g, root = "00")

heuristicAND(S, g, root = "00")

heuristicOR(S, g, root = "00")
```

Arguments

S	a named flat scores matrix with examples on rows and classes on columns
g	a graph of class graphNEL. It represents the hierarchy of the classes
root	name of the class that it is the top-level (root) of the hierarchy (def:00)

Details

Heuristic Methods:

1. **MAX**: reports the largest logist regression (LR) value of self and all descendants: $p_i = max_{j \in descendants(i)} \hat{p_j}$;
2. **AND**: reports the product of LR values of all ancestors and self. This is equivalent to computing the probability that all ancestral terms are "on" assuming that, conditional on the data, all predictions are independent: $p_i = \prod_{j \in ancestors(i)} \hat{p_j}$;
3. **OR**: computes the probability that at least one of the descendant terms is "on" assuming again that, conditional on the data, all predictions are independent: $1 - p_i = \prod_{j \in descendants(i)} (1 - \hat{p_j})$;

Value

a matrix with the scores of the classes corrected according to the chosen heuristic algorithm

Examples

```
data(graph);
data(scores);
data(labels);
root <- root.node(g);
S.heuristicMAX <- heuristicMAX(S,g,root);
S.heuristicAND <- heuristicAND(S,g,root);
S.heuristicOR <- heuristicOR(S,g,root);
```

---

hierarchical.checkers *Hierarchical constraints cheker*

Description

Check if the true path rule is violated or not. In other words this function checks if the score of a parent or an ancestor node is always larger or equal than that of its children or descendants nodes

Usage

```
check.hierarchy.single.sample(y.hier, g, root = "00")

check.hierarchy(S.hier, g, root = "00")
```

Arguments

y.hier	vector of scores relative to a single example. This must be a named numeric vector
g	a graph of class `graphNEL`. It represents the hierarchy of the classes
root	name of the class that it is the top-level (root) of the hierarchy (def:00)
S.hier	the matrix with the scores of the classes corrected in according to hierarchy. This must be a named matrix: rows are examples and colums are classes

Value

return a list of 3 elements:

- Status:
    - OK if none hierarchical constraints have bee broken;
    - NOTOK if there is at least one hierarchical constraints broken;
- Hierarchy_Constraints_Broken:
    - TRUE: example did not respect the hierarchical constraints;
    - FALSE: example broke the hierarchical constraints;
- Hierarchy_costraints_satisfied: how many terms satisfied the hierarchical constraint

Examples

```
data(graph);
data(scores);
root <- root.node(g);
S.hier <- htd(S,g,root);
S.hier.single.example <- S.hier[sample(ncol(S.hier),1),];
check.hierarchy.single.sample(S.hier.single.example, g, root=root);
check.hierarchy(S.hier, g, root);
```

HTD-DAG   *HTD-DAG*

Description

Implementation of a top-down procedure to correct the scores of the hierarchy according to the constraints that the score of a node cannot be greater than a score of its parents.

Usage

```
htd(S, g, root = "00")
```

Arguments

S	a named flat scores matrix with examples on rows and classes on columns
g	a graph of class `graphNEL`. It represents the hierarchy of the classes
root	name of the class that it is the top-level (root) of the hierarchy (def:00)

Details

The HTD-DAG algorithm modifies the flat scores according to the hierarchy of a DAG through a unique run across the nodes of the graph. For a given example $x \in X$, the flat predictions $f(x) = \hat{y}$ are hierarchically corrected to $\bar{y}$, by per-level visiting the nodes of the DAG from top to bottom according to the following simple rule:

$$\bar{y}_i := \begin{cases} \hat{y}_i & \text{if} \quad i \in root(G) \\ \min_{j \in par(i)} \bar{y}_j & \text{if} \quad \min_{j \in par(i)} \bar{y}_j < \hat{y}_i \\ \hat{y}_i & \text{otherwise} \end{cases}$$

The node levels correspond to their maximum path length from the root.

Value

a matrix with the scores of the classes corrected according to the HTD-DAG algorithm.

See Also

graph.levels, hierarchical.checkers

Examples

```
data(graph);
data(scores);
root <- root.node(g);
S.htd <- htd(S,g,root);
```

Multilabel.F.measure   *Multilabel F-measure*

Description

Method for computing Precision, Recall, Specificity, Accuracy and F-measure for multiclass and multilabel classification

Usage

```
F.measure.multilabel(target, predicted, b.per.example = FALSE)

## S4 method for signature 'matrix,matrix'
F.measure.multilabel(target, predicted,
  b.per.example = FALSE)
```

Arguments

target   matrix with the target multilabels: rows correspond to examples and columns to classes. $target[i, j] = 1$ if example i belongs to class j, $target[i, j] = 0$ otherwise

predicted   a numeric matrix with discrete predicted values: rows correspond to examples and columns to classes. $predicted[i, j] = 1$ if example i is predicted belonging to class j, $target[i, j] = 0$ otherwise

b.per.example   boolean.

- TRUE: results are returned for each example;
- FALSE: only the average results are returned

Details

Names of rows and columns of target and predicted matrix must be provided in the same order, otherwise a stop message is returned

Value

Two different outputs respect to the input parameter b.per.example:

- b.per.example==FALSE: a list with a single element average. A named vector with average precision (P), recall (R), specificity (S), F-measure (F), average F-measure (avF) and Accuracy (A) across examples. F is the F-measure computed as the harmonic mean between the average precision and recall; av.F is the F-measure computed as the average across examples.

- `b.per.example==FALSE`: a list with two elements:

    1. average: a named vector with average precision (P), recall (R), specificity (S), F-measure (F), average F-measure (avF) and Accuracy (A) across examples;
    2. per.example: a named matrix with the Precision (P), Recall (R), Specificity (S), Accuracy (A), F-measure (F) and av.F-measure (av.F) for each example. Row names correspond to examples, column names correspond respectively to Precision (P), Recall (R), Specificity (S), Accuracy (A), F-measure (F) and av.F-measure (av.F)

Examples

```
data(labels);
data(scores);
data(graph);
root <- root.node(g);
L <- L[,-which(colnames(L)==root)];
S <- S[,-which(colnames(S)==root)];
S[S>0.7] <- 1;
S[S<0.7] <- 0;
FMM <- F.measure.multilabel(L,S);
```

---

normalize.max	*Max normalization*

Description

Function to normalize the scores of a flat scores matrix per class

Usage

```
normalize.max(S)
```

Arguments

S   matrix with the raw non normalized scores. Rows are examples and columns are classes

Details

The scores of each class are normalized by dividing the score values for the maximum score of that class. If the max score of a class is zero, no normalization is needed, otherwise NaN value will be printed as results of 0 out of 0 division.

Value

A score matrix with the same dimensions of S, but with scores max/normalized separately for each class

Examples

```
data(scores);
maxnorm <- normalize.max(S);
```

---

parents   *Build parents*

Description

Compute the parents for each node of a graph

Usage

```
get.parents(g, root = "00")

get.parents.top.down(g, levels, root = "00")

get.parents.bottom.up(g, levels, root = "00")

get.parents.topological.sorting(g, root = "00")
```

Arguments

g	a graph of class `graphNEL`. It represents the hierarchy of the classes
root	name of the root node (def. `root="00"`)
levels	a list of character vectors. Each component represents a graph level and the elements of any component correspond to nodes. The level 0 coincides with the root node.

Value

`get.parents` returns a named list of character vectors. Each component corresponds to a node x of the graph (i.e. child node) and its vector is the set of its parents (the root node is not included)

`get.parents.top.down` returns a named list of character vectors. Each component corresponds to a node x of the graph (i.e. child node) and its vector is the set of its parents. The nodes order follows the levels of the graph from root (excluded) to leaves.

`get.parents.bottom.up` returns a named list of character vectors. Each component corresponds to a node x of the graph (i.e. child node) and its vector isthe set of its parents. The nodes are ordered from leaves to root (excluded).

`get.parents.topological.sorting` a named list of character vectors. Each component corresponds to a node x of the graph (i.e. child node) and its vector is the set of its parents. The nodes are ordered according to a topological sorting, i.e. parents node come before children node.

See Also

[graph.levels](#)

Examples

```
data(graph);
root <- root.node(g)
parents <- get.parents(g, root=root);
lev <- graph.levels(g, root=root);
parents.tod <- get.parents.top.down(g, lev, root=root);
parents.bup <- get.parents.bottom.up(g, lev, root=root);
parents.tsort <- get.parents.topological.sorting(g, root=root);
```

PXR   *Precision at fixed Recall level*

Description

Function to compute the Precision at fixed Recall levels (PXR) over classes through **PerfMeas** package

Usage

```
PXR.at.multiple.recall.levels.over.classes(target, pred, rec.levels = seq(from
  = 0.1, to = 1, by = 0.1), folds = NULL, seed = NULL)
```

Arguments

target   matrix with the target multilabels: rows correspond to examples and columns to classes. $target[i, j] = 1$ if example i belongs to class j, $target[i, j] = 0$ otherwise.

pred   a numeric matrix with predicted values (scores): rows correspond to examples and columns to classes.

rec.levels   a vector with the desired recall levels (def: from:0.1, to:0.9, by:0.1)

folds   number of folds on which computing the AUROC. If folds=NULL (def.), the one-shot AUROC is computed, otherwise the AUROC is computed averaged across folds.

seed   initialization seed for the random generator to create folds. Set seed only if folds≠NULL. If seed≠NULL and folds=NULL, the AUROC averaged across folds is computed without seed initialization.

Details

The PXR across all classes is computed either one-shot or averaged across stratified folds.

Function to compute the precision at fixed recall levels across all classes either one-shot or averaged across stratified folds

Names of rows and columns of target and pred matrix must be provided in the same order, otherwise a stop message is returned

Value

a list with two elements:

1. average: a vector with the the average precisions at different recall levels across classes

2. per.class: a matrix with the precisions at different recall levels: rows are classes, columns precisions at different recall level

Examples

```
data(labels);
data(scores);
data(graph);
root <- root.node(g);
L <- L[,-which(colnames(L)==root)];
S <- S[,-which(colnames(S)==root)];
rec.levels <- seq(from=0.1, to=1, by=0.1);
PXR <- PXR.at.multiple.recall.levels.over.classes(L, S, rec.levels=rec.levels, folds=5, seed=23);
```

---

read.graph   *Read a directed graph from a file*

Description

A directed graph is read from a file and a `graphNEL` object is built

Usage

```
read.graph(file = "graph.txt")
```

Arguments

file   name of the file to be read. The format of the file is a sequence of rows and each
                  row corresponds to an edge represented through a pair of vertices separated by
                  blanks

Value

an object of class `graphNEL`

Examples

```
ed <- system.file("extdata/graph.edges.txt", package= "HEMDAG");
g <- read.graph(file=ed);
```

---

read.undirected.graph   *Read an undirected graph from a file*

Description

The graph is read from a file and a `graphNEL` object is built. The format of the input file is a sequence of rows. Each row corresponds to an edge represented through a pair of vertices separated by blanks, and the weight of the edge.

Usage

```
read.undirected.graph(file = "graph.txt")
```

Arguments

file   name of the file to be read

Value

a graph of class `graphNEL`

Examples

```
edges <- system.file("extdata/edges.txt" ,package="HEMDAG");
g <- read.undirected.graph(file=edges);
```

---

root.node   *Root node*

Description

Find the root node of a directed graph

Usage

```
root.node(g)
```

Arguments

g   a graph of class `graphNEL`. It represents the hierarchy of the classes

Value

name of the root node

Examples

```
data(graph);
root <- root.node(g);
```

---

scores.normalization *Scores Normalization Function*

Description

Functions to normalize a flat scores matrix w.r.t. max normalization (MaxNorm) or quantile normalization (Qnorm)

Usage

```
scores.normalization(norm.type = "MaxNorm", S)
```

Arguments

norm.type
: can be one of the following two values:

  - MaxNorm (def.): each score is divided w.r.t. the max of each class;
  - Qnorm: a quantile normalization is applied. Library preprocessCore is used.

S
: a named flat scores matrix with examples on rows and classes on columns

Details

To apply the quantile normalization the **preprocessCore** library is uded.

Value

the matrix of the scores flat normalized w.r.t. MaxNorm or Qnorm

Examples

```
data(scores);
norm.types <- c("MaxNorm","Qnorm");
for(norm.type in norm.types){
scores.normalization(norm.type=norm.type, S=S)
}
```

---

specific.annotation.list
                    Specific annotations list

Description

Construct a list of the most specific annotations starting from the table of the most specific annotations

Usage

```
specific.annotation.list(ann)
```

Arguments

ann   annotation matrix (0/1). Rows are examples and columns are most specific
                       terms. It must be a named matrix.

Value

a named list, where the names of each component correspond to an examples (genes) and the ele-
ments of each component are the most specific classes associated to that genes

See Also

[specific.annotation.matrix](#)

Examples

```
data(labels);
spec.list <- specific.annotation.list(L);
```

---

specific.annotation.matrix
                          HPO specific annotations matrix

Description

Construct the labels matrix of the most specific HPO terms

Usage

```
specific.annotation.matrix(file = "gene2pheno.txt", genename = "TRUE")
```

Arguments

file   text file representing the most specific associations gene-HPO term
                       (def: "gene2pheno.txt"). The file must be written as sequence of rows. Each
                       row represents a gene and all its associations with abnormal phenotype tab sep-
                       arated,
                       e.g.: gene_1 <tab> phen1 <tab> ... phen_N.
                       See **Details** section to know more information about how to obtain this file.

genename   boolean value:

                            • TRUE (def.): the names of genes are *gene symbol* (i.e. characters);
                            • FALSE: the names of gene are entrez *gene ID* (i.e. integer numbers);

Details

The input plain text file representing the most specific associations gene-HPO term can be obtained by forking the GitHub repositoty HPOparser, a collection of `Perl` subroutines to parse the HPO OBO file and the HPO annotations file.

Value

the annotation matrix of the most specific annotations (0/1): rows are genes and columns are HPO terms. Let's denote M the labels matrix. If $M[i, j] = 1$, means that the gene i is annotated with the class j, otherwise $M[i, j] = 0$.

Examples

```
gene2pheno <- system.file("extdata/gene2pheno.txt", package="HEMDAG");
spec.ann <- specific.annotation.matrix(gene2pheno, genename=TRUE);
```

---

stratified.cross.validation

Stratified cross validation

Description

Generate data for the stratified cross-validation

Usage

```
do.stratified.cv.data.single.class(examples, positives, kk = 5, seed = NULL)

do.stratified.cv.data.over.classes(labels, examples, kk = 5, seed = NULL)
```

Arguments

examples	indices or names of the examples. Can be either a vector of integers or a vector of names.
positives	vector of integers or vector of names. The indices (or names) refer to the indices (or names) of 'positive' examples
kk	number of folds (def=5)
seed	seed of the random generator (def=NULL). If is set to NULL no initialization is performed
labels	labels matrix. Rows are genes and columns are classes. Let's denote M the labels matrix. If $M[i, j] = 1$, means that the gene i is annotated with the class j, otherwise $M[i, j] = 0$.

Value

do.stratified.cv.data.single.class returns a list with 2 two component:

- fold.non.positives: a list with k components. Each component is a vector with the indices (or names) of the non-positive elements. Indices (or names) refer to row numbers (or names) of a data matrix.

- fold.positives: a list with k components. Each component is a vector with the indices (or names) of the positive elements. Indices (or names) refer to row numbers (or names) of a data matrix.

do.stratified.cv.data.over.classes returns a list with n components, where n is the number of classes of the labels matrix. Each component n is in turn a list with k elements, where k is the number of folds. Each fold contains an equal amount of examples positives and negatives.

Examples

```
data(labels);
examples.index <- 1:nrow(L);
examples.name <- rownames(L);
positives <- which(L[,3]==1);
x <- do.stratified.cv.data.single.class(examples.index, positives, kk=5, seed=23);
y <- do.stratified.cv.data.single.class(examples.name, positives, kk=5, seed=23);
z <- do.stratified.cv.data.over.classes(L, examples.index, kk=5, seed=23);
k <- do.stratified.cv.data.over.classes(L, examples.name, kk=5, seed=23);
```

TPR-DAG   *TPR-DAG variants*

Description

Different variants of the TPR-DAG algorithm are implemented. In their more general form the TPR-DAG algorithms adopt a two step learnig strategy:

1. in the first step they compute a *per-level bottom-up* visit from the leaves to the root to propagate positive predictions across the hierarchy;

2. in the second step they compute a *per-level top-down* visit from the root to the leaves in order to assure the hierarchical consistency of the predictions

Usage

```
tpr.threshold(S, g, root = "00", t = 0.5)

tpr.threshold.free(S, g, root = "00")

tpr.weighted.threshold.free(S, g, root = "00", w = 0.5)

tpr.weighted.threshold(S, g, root = "00", t = 0.5, w = 0.5)
```

Arguments

S	a named flat scores matrix with examples on rows and classes on columns
g	a graph of class `graphNEL`. It represents the hierarchy of the classes
root	name of the class that it is on the top-level of the hierarchy (def. `root="00"`)
t	threshold for the choice of positive children (def. `t=0.5`)
w	weight to balance between the contribution of the node i and that of its positive children

Details

The *vanilla* TPR-DAG adopts a per-level bottom-up traversal of the DAG to correct the flat predictions $\hat{y}_i$:

$$\bar{y}_i := \frac{1}{1 + |\phi_i|} (\hat{y}_i + \sum_{j \in \phi_i} \bar{y}_j)$$

where ϕ_i are the positive children of i. Different strategies to select the positive children ϕ_i can be applied:

1. **Threshold-Free** strategy: the positive nodes are those children that can increment the score of the node i, that is those nodes that achieve a score higher than that of their parents:

$$\phi_i := \{j \in child(i) | \bar{y}_j > \hat{y}_i\}$$

2. **Threshold** strategy: the positive children are selected on the basis of a threshold that can ben selected in two different ways:

   (a) for each node a constant threshold $\bar{t}$ is a priori selected:

$$\phi_i := \{j \in child(i) | \bar{y}_j > \bar{t}\}$$

   For instance if the predictions represent probabilities it could be meaningful to a priori select $\bar{t} = 0.5$.

   (b) the threshold is selected to maximize some performance metric $\mathcal{M}$ estimated on the training data, as for instance the F-score or the AUPRC. In other words the threshold is selected to maximize some measure of accuracy of the predictions $\mathcal{M}(j, t)$ on the training data for the class j with respect to the threshold t. The corresponding set of positives $\forall i \in V$ is:

$$\phi_i := \{j \in child(i) | \bar{y}_j > t_j^*, t_j^* = \arg\max_t \mathcal{M}(j, t)\}$$

   For instance t_j^* can be selected from a set of $t \in (0, 1)$ through internal cross-validation techniques.

The weighted TPR-DAG version can be designed by adding a weight $w \in [0, 1]$ to balance between the contribution of the node i and that of its positive children ϕ, through their convex combination:

$$\bar{y}_i := w\hat{y}_i + \frac{(1 - w)}{|\phi_i|} \sum_{j \in \phi_i} \bar{y}_j$$

If $w = 1$ no weight is attributed to the children and the TPR-DAG reduces to the HTD-DAG algorithm, since in this way only the prediction for node i is used in the bottom-up step of the algorithm. If $w = 0$ only the predictors associated to the children nodes vote to predict node i. In the intermediate cases we attribute more importance to the predictor for the node i or to its children depending on the values of w.

Value

a named matrix with the scores of the classes corrected according to the TPR-DAG algorithm.

Examples

```
data(graph);
data(scores);
data(labels);
root <- root.node(g);
S.tprTF <- tpr.threshold.free(S,g,root);
S.tprT <- tpr.threshold(S,g,root,t=0.5);
S.tprW <- tpr.weighted.threshold.free(S,g,root,w=0.5);
S.tprWT <- tpr.weighted.threshold(S,g,root,w=0.5, t=0.5);
```

TPR-DAG-cross-validation

TPR-DAG cross-validation experiments

Description

High level function to correct the computed scores in a hierarchy according to the chosen ensemble algorithm

Usage

```
Do.TPR.DAG(threshold = seq(from = 0.1, to = 0.9, by = 0.1),
  weight = seq(from = 0.1, to = 0.9, by = 0.1), kk = 5, folds = 5,
  seed = 23, norm = TRUE, norm.type = NULL, positive = "children",
  bottomup = "threshold.free", rec.levels = seq(from = 0.1, to = 1, by =
  0.1), n.round = 3, f.criterion = "F", metric = NULL,
  flat.file = flat.file, ann.file = ann.file, dag.file = dag.file,
  flat.dir = flat.dir, ann.dir = ann.dir, dag.dir = dag.dir,
  hierScore.dir = hierScore.dir, perf.dir = perf.dir)
```

Arguments

threshold	range of threshold values to be tested in order to find the best threshold (def: from:0.1, to:0.9, by:0.1). The denser the range is, the higher the probability to find the best threshold is, but obviously the execution time will be higher. Set the parameter threshold only for the variants that requiring a threshold for the positive nodes selection, otherwise set the parameter threshold to zero
weight	range of weight values to be tested in order to find the best weight (def: from:0.1, to:0.9, by:0.1). The denser the range is, the higher the probability to find the best threshold is, but obviously the execution time will be higher. Set the parameter weight only for the *weighted* variants, otherwise set the parameter weight to zero

kk	number of folds of the cross validation (def: kk=5) on which tuning the parameters threshold and weight and
folds	number of folds of the cross validation on which computing the performance metrics averaged across folds (def. 5). If folds=NULL, the performance metrics are computed one-shot, otherwise the performance metrics are averaged across folds.
seed	initialization seed for the random generator to create folds (def. 23). If NULL folds are generated without seed initialization. The parameter seed controls both the parameter kk and the parameter folds.
norm	boolean value: should the flat scores matrix be normalized?

- TRUE (def.): the flat scores matrix has been already normalized in according to a normalization method;
- FALSE: the flat scores matrix has not been normalized yet. See the parameter norm.type to set the on the fly normalization method to apply among those possible.

norm.type	can be one of the following three values:

1. NULL (def.): set norm.type to NULL if and only if the parameter norm is set to TRUE;
2. MaxNorm: each score is divided for the maximum of each class;
3. Qnorm: quantile normalization. **preprocessCore** package is used.

positive	choice of the *positive* nodes to be considered in the bottom-up strategy. Can be one of the following values:

- children (def.): for each node are considered its positive children;
- descendants: for each node are considered its positive descendants;

bottomup	strategy to enhance the flat predictions by propagating the positive predictions from leaves to root. It can be one of the following values:

- threshold.free (def.): positive nodes are selected on the basis of the threshold.free strategy (def.);
- threshold: positive nodes are selected on the basis of the threshold strategy;
- weighted.threshold.free: positive nodes are selected on the basis of the weighted.threshold.free strategy;
- weighted.threshold: positive nodes are selected on the basis of the weighted.threshold strategy;
- tau: positive nodes are selected on the basis of the tau strategy. NOTE: tau is only a DESCENS variants. If you use tau strategy you must set the parameter positive=descendants;

rec.levels	a vector with the desired recall levels (def: from:0.1, to:0.9, by:0.1) to compute the the Precision at fixed Recall level (PXR)
n.round	number of rounding digits to be applied to the hierarchical scores matrix (def. 3). It is used for choosing the best threshold on the basis of the best F-measure
f.criterion	character. Type of F-measure to be used to select the best F-measure. Two possibilities:

1. `F` (def.): corresponds to the harmonic mean between the average precision and recall

2. `avF`: corresponds to the per-example F-score averaged across all the examples

metric   a string character specifying the performance metric on which to maximize the parametric ensemble variant. It can be one of the following values:

1. `PRC`: the parametric ensemble variant is maximized on the basis of AUPRC ([AUPRC](#));

2. `FMAX`: the parametric ensemble variant is maximized on the basis of Fmax ([Multilabel.F.measure](#);

3. `NULL`: on the `threshold.free` variant none parameter optimization is needed, since the variant is non-parametric. So, if `bottomup=threshold.free` set `metric=NULL` (def.).

flat.file   name of the file containing the flat scores matrix to be normalized or already normalized (without rda extension)

ann.file   name of the file containing the the label matrix of the examples (without rda extension)

dag.file   name of the file containing the graph that represents the hierarchy of the classes (without rda extension)

flat.dir   relative path where flat scores matrix is stored

ann.dir   relative path where annotation matrix is stored

dag.dir   relative path where graph is stored

hierScore.dir   relative path where the hierarchical scores matrix must be stored

perf.dir   relative path where the performance measures must be stored

Details

The parametric hierarchical ensemble variants are cross-validated by maximizing in according to the metric chosen in the parameter `metric`, that is F-measure ([Multilabel.F.measure](#)) or AUPRC ([AUPRC](#)).

The function checks if the number of classes between the flat scores matrix and the annotations matrix mismatched. If so, the number of terms of the annotations matrix is shrunk to the number of terms of the flat scores matrix and the corresponding subgraph is computed as well. N.B.: it is supposed that all the nodes of the subgraph are accessible from the root.

Value

Two `rda` files stored in the respective output directories:

1. `Hierarchical Scores Results`: a matrix with examples on rows and classes on columns representing the computed hierarchical scores for each example and for each considered class. It is stored in the `hierScore.dir` directory.

2. `Performance Measures`: *flat* and *hierarchical* performace results:

    (a) AUPRC results computed though `AUPRC.single.over.classes` ([AUPRC](#));

    (b) AUROC results computed through `AUROC.single.over.classes` ([AUROC](#));

(c) PXR results computed though PXR.at.multiple.recall.levels.over.classes (PXR);

(d) FMM results computed though compute.Fmeasure.multilabel (FMM);

It is stored in the perf.dir directory.

See Also

TPR-DAG-variants

Examples

```
data(graph);
data(scores);
data(labels);
if (!dir.exists("data")){
dir.create("data");
}
if (!dir.exists("results")){
dir.create("results");
}
save(g,file="data/graph.rda");
save(L,file="data/labels.rda");
save(S,file="data/scores.rda");
dag.dir <- flat.dir <- ann.dir <- "data/";
hierScore.dir <- perf.dir <- "results/";
dag.file <- "graph";
flat.file <- "scores";
ann.file <- "labels";
threshold <- weight <- 0;
norm.type <- "MaxNorm";
positive <- "children";
bottomup <- "threshold.free";
rec.levels <- seq(from=0.1, to=1, by=0.1);
Do.TPR.DAG(threshold=threshold, weight=weight, kk=5, folds=5, seed=23, norm=FALSE,
norm.type=norm.type, positive=positive, bottomup=bottomup, n.round=3, f.criterion="F",
metric=NULL, rec.levels=rec.levels, flat.file=flat.file, ann.file=ann.file,
dag.file=dag.file, flat.dir=flat.dir, ann.dir=ann.dir, dag.dir=dag.dir,
hierScore.dir=hierScore.dir, perf.dir=perf.dir);
```

TPR-DAG-holdout   *TPR-DAG holdout experiments*

Description

High level function to correct the computed scores in a hierarchy according to the chosen ensemble algorithm

Usage

```
Do.TPR.DAG.holdout(threshold = seq(from = 0.1, to = 0.9, by = 0.1),
  weight = seq(from = 0.1, to = 1, by = 0.1), kk = 5, folds = 5,
  seed = 23, norm = TRUE, norm.type = NULL, positive = "children",
  bottomup = "threshold.free", rec.levels = seq(from = 0.1, to = 1, by =
  0.1), n.round = 3, f.criterion = "F", metric = NULL,
  flat.file = flat.file, ann.file = ann.file, dag.file = dag.file,
  ind.test.set = ind.test.set, ind.dir = ind.dir, flat.dir = flat.dir,
  ann.dir = ann.dir, dag.dir = dag.dir, hierScore.dir = hierScore.dir,
  perf.dir = perf.dir)
```

Arguments

threshold	range of threshold values to be tested in order to find the best threshold (def: `from:0.1, to:0.9, by:0.1`). The denser the range is, the higher the probability to find the best threshold is, but obviously the execution time will be higher. Set the parameter `threshold` only for the variants that requiring a threshold for the positive nodes selection, otherwise set the parameter `threshold` to zero
weight	range of weight values to be tested in order to find the best weight (def: `from:0.1, to:0.9, by:0.1`). The denser the range is, the higher the probability to find the best threshold is, but obviously the execution time will be higher. Set the parameter `weight` only for the *weighted* variants, otherwise set the parameter `weight` to zero
kk	number of folds of the cross validation (def: `kk=5`) on which tuning the parameters `threshold` and `weight` and
folds	number of folds of the cross validation on which computing the performance metrics averaged across folds (def. 5). If `folds=NULL`, the performance metrics are computed one-shot, otherwise the performance metrics are averaged across folds.
seed	initialization seed for the random generator to create folds (def. 23). If `NULL` folds are generated without seed initialization. The parameter seed controls both the parameter kk and the parameter `folds`.
norm	boolean value: should the flat scores matrix be normalized? • `TRUE` (def.): the flat scores matrix has been already normalized in according to a normalization method; • `FALSE`: the flat scores matrix has not been normalized yet. See the parameter `norm.type` to set the on the fly normalization method to apply among those possible.
norm.type	can be one of the following three values: 1. `NULL` (def.): set `norm.type` to `NULL` if and only if the parameter `norm` is set to `TRUE`; 2. `MaxNorm`: each score is divided for the maximum of each class; 3. `Qnorm`: quantile normalization. **preprocessCore** package is used.
positive	choice of the *positive* nodes to be considered in the bottom-up strategy. Can be one of the following values:

- `children` (def.): for each node are considered its positive children;
- `descendants`: for each node are considered its positive descendants;

bottomup   strategy to enhance the flat predictions by propagating the positive predictions from leaves to root. It can be one of the following values:

- `threshold.free` (def.): positive nodes are selected on the basis of the `threshold.free` strategy (def.);
- `threshold`: positive nodes are selected on the basis of the `threshold` strategy;
- `weighted.threshold.free`: positive nodes are selected on the basis of the `weighted.threshold.free` strategy;
- `weighted.threshold`: positive nodes are selected on the basis of the `weighted.threshold` strategy;
- `tau`: positive nodes are selected on the basis of the `tau` strategy. NOTE: `tau` is only a `DESCENS` variants. If you use `tau` strategy you must set the parameter `positive=descendants`;

rec.levels   a vector with the desired recall levels (def: `from:0.1`, `to:0.9`, `by:0.1`) to compute the the Precision at fixed Recall level (PXR)

n.round   number of rounding digits to be applied to the hierarchical scores matrix (def. 3). It is used for choosing the best threshold on the basis of the best F-measure

f.criterion   character. Type of F-measure to be used to select the best F-measure. Two possibilities:

1. `F` (def.): corresponds to the harmonic mean between the average precision and recall
2. `avF`: corresponds to the per-example F-score averaged across all the examples

metric   a string character specifying the performance metric on which to maximize the parametric ensemble variant. It can be one of the following values:

1. `PRC`: the parametric ensemble variant is maximized on the basis of AUPRC ([AUPRC](#));
2. `FMAX`: the parametric ensemble variant is maximized on the basis of Fmax ([Multilabel.F.measure](#);
3. `NULL`: on the `threshold.free` variant none parameter optimization is needed, since the variant is non-parametric. So, if `bottomup=threshold.free` set `metric=NULL` (def.).

flat.file   name of the file containing the flat scores matrix to be normalized or already normalized (without rda extension)

ann.file   name of the file containing the the label matrix of the examples (without rda extension)

dag.file   name of the file containing the graph that represents the hierarchy of the classes (without rda extension)

ind.test.set   name of the file containing a vector of integer numbers corresponding to the indices of the elements (rows) of scores matrix to be used in the test set

ind.dir   relative path to folder where `ind.test.set` is stored

flat.dir	relative path where flat scores matrix is stored
ann.dir	relative path where annotation matrix is stored
dag.dir	relative path where graph is stored
hierScore.dir	relative path where the hierarchical scores matrix must be stored
perf.dir	relative path where the performance measures must be stored

Details

The parametric hierarchical ensemble variants are cross-validated by maximizing in according to the metric chosen in the parameter metric, that is F-measure (Multilabel.F.measure) or AUPRC (AUPRC).

The function checks if the number of classes between the flat scores matrix and the annotations matrix mismatched. If so, the number of terms of the annotations matrix is shrunk to the number of terms of the flat scores matrix and the corresponding subgraph is computed as well. N.B.: it is supposed that all the nodes of the subgraph are accessible from the root.

Value

Two rda files stored in the respective output directories:

1. Hierarchical Scores Results: a matrix with examples on rows and classes on columns representing the computed hierarchical scores for each example and for each considered class. It is stored in the hierScore.dir directory.

2. Performance Measures: *flat* and *hierarchical* performace results:
   (a) AUPRC results computed though AUPRC.single.over.classes (AUPRC);
   (b) AUROC results computed through AUROC.single.over.classes (AUROC);
   (c) PXR results computed though PXR.at.multiple.recall.levels.over.classes (PXR);
   (d) FMM results computed though compute.Fmeasure.multilabel (FMM);

It is stored in the perf.dir directory.

See Also

TPR-DAG-variants

Examples

```
data(graph);
data(scores);
data(labels);
data(test.index);
if (!dir.exists("data")){
dir.create("data");
}
if (!dir.exists("results")){
dir.create("results");
}
save(g,file="data/graph.rda");
```

```
save(L,file="data/labels.rda");
save(S,file="data/scores.rda");
save(test.index, file="data/test.index.rda");
ind.dir <- dag.dir <- flat.dir <- ann.dir <- "data/";
hierScore.dir <- perf.dir <- "results/";
dag.dir <- flat.dir <- ann.dir <- "data/";
ind.test.set <- "test.index";
dag.file <- "graph";
flat.file <- "scores";
ann.file <- "labels";
threshold <- weight <- 0;
norm.type <- "MaxNorm";
positive <- "children";
bottomup <- "threshold.free";
rec.levels <- seq(from=0.1, to=1, by=0.1);
Do.TPR.DAG.holdout(threshold=threshold, weight=weight, kk=5, folds=NULL, seed=23, norm=FALSE,
norm.type=norm.type, positive=positive, bottomup=bottomup, rec.levels=rec.levels, n.round=3,
f.criterion="F", metric=NULL, flat.file=flat.file, ann.file=ann.file, dag.file=dag.file,
ind.test.set=ind.test.set, ind.dir=ind.dir, flat.dir=flat.dir, ann.dir=ann.dir,
dag.dir=dag.dir, hierScore.dir=hierScore.dir, perf.dir=perf.dir);
```

TPR-DAG-variants   *TPR-DAG Variants*

Description

Function gathering several hierarchical ensemble algorithms

Usage

```
TPR.DAG(S, g, root = "00", positive = "children",
  bottomup = "threshold.free", t = 0, w = 0)
```

Arguments

S	a named flat scores matrix with examples on rows and classes on columns
g	a graph of class `graphNEL`. It represents the hierarchy of the classes
root	name of the class that it is on the top-level of the hierarchy (def. root="00")
positive	choice of the *positive* nodes to be considered in the bottom-up strategy. Can be one of the following values:

- children (def.): for each node are considered its positive children;
- descendants: for each node are considered its positive descendants;

bottomup	strategy to enhance the flat predictions by propagating the positive predictions from leaves to root. It can be one of the following values:

- threshold.free (def.): positive nodes are selected on the basis of the threshold.free strategy (def.);

- threshold: positive nodes are selected on the basis of the threshold strategy;
- weighted.threshold.free: positive nodes are selected on the basis of the weighted.threshold.free strategy;
- weighted.threshold: positive nodes are selected on the basis of the weighted.threshold strategy;
- tau: positive nodes are selected on the basis of the tau strategy. NOTE: tau is only a DESCENS variants. If you use tau strategy you must set the parameter positive=descendants;

t   threshold for the choice of positive nodes (def. t=0.5). Set t only for the variants that requiring a threshold for the selection of the positive nodes, otherwise set t to zero

w   weight to balance between the contribution of the node i and that of its positive nodes. Set w only for the *weighted* variants, otherwise set w to zero

Value

a named matrix with the scores of the classes corrected according to the chosen algorithm

See Also

[TPR-DAG](#), [DESCENS](#), [HTD-DAG](#)

Examples

```
data(graph);
data(scores);
data(labels);
root <- root.node(g);
S.hier <- TPR.DAG(S, g, root, positive="children", bottomup="threshold.free", t=0, w=0);
```

---

transitive.closure.annotations
                    Transitive closure of annotations

Description

Performs the transitive closure of the annotations using ancestors and the most specific annotation table. The annotations are propagated from bottom to top, enriching the most specific annotations table. The rows of the matrix correspond to the genes of the most specific annotation table and the columns to the HPO terms/classes

Usage

```
transitive.closure.annotations(ann.spec, anc)
```

Arguments

ann.spec	the annotation matrix of the most specific annotations (0/1): rows are genes and columns are HPO terms.
anc	list of the ancestors of the ontology.

Value

an annotation table T: rows correspond to genes and columns to HPO terms. $T[i, j] = 1$ means that gene i is annotated for the term j, $T[i, j] = 0$ means that gene i is not annotated for the term j.

See Also

[specific.annotation.matrix](), [build.ancestors]()

Examples

```
data(graph);
data(labels);
anc <- build.ancestors(g);
tca <- transitive.closure.annotations(L, anc);
```

---

tupla.matrix *Tupla Matrix*

Description

Trasform a Weighted Adjacency Matrix (wadj matrix) of a graph in a tupla, i.e. as a sequences of rows separated by blank and the weight of the edges, e.g nodeX nodeY score

Usage

```
tupla.matrix(m, compressed = TRUE, output.file = "net.file.gz")
```

Arguments

m	a weighetd adjacency matrix of the graph. Rows and columns are examples. It must be a square named matrix.
compressed	boolean value:
	• TRUE (def.): the output file will be a gz compressed format;
	• FALSE: the output file will be a plain text format;
output.file	name of the file of the to be written

Details

Only the *non-zero* interactions are kept, while the *zero* interactions are discarded. In other words in the output.file are reported only those nodes having a weight different from zero

Value

if `compressed=TRUE` the weighted adjacency matrix as tupla is stored in a compressed gz, otherwise (`compressed=FALSE`) it is stored in a plain `text` file.

Examples

```
## Not run:
data(wadj);
tupla.matrix(W,compressed=TRUE, output.file="tupla.wadj.gz");
tupla.matrix(W,compressed=FALSE, output.file="tupla.wadj.txt");
## End(Not run)
```

---

weighted.adjacency.matrix

Weighted Adjacency Matrix

Description

Construct a Weighted Adjacency Matrix (wadj matrix) of a graph

Usage

```
weighted.adjacency.matrix(file = "edges.txt", compressed = TRUE,
  nodename = TRUE)
```

Arguments

`file`	name of the plain text file to be read (def. edges). The format of the file is a sequence of rows. Each row corresponds to an edge represented through a pair of vertices separated by blanks and the weight of the edges. For instance: `nodeX nodeY score`
`compressed`	boolean value: • TRUE (def.): the input file must be in a .gz compressed format; • FALSE: the input file must be in a plain text format;
`nodename`	boolean value: • TRUE (def.): the names of nodes are gene symbol (i.e. characters); • FALSE: the names of the nodes are entrez gene ID (i.e. integer numbers);

Details

The input parameter nodename sorts the row names of the wadj matrix in increasing order if they are integer number or in alphabetic order if they are characters.

Value

a named symmetric weighted adjacency matrix of the graph

Examples

```
edges <- system.file("extdata/edges.txt", package="HEMDAG");
W <- weighted.adjacency.matrix(file=edges, compressed=FALSE, nodename=TRUE);
```

---

write.graph   *Write a directed graph on file*

Description

An object of class graphNEL is read and the graph is written on a plain text file as sequence of rows

Usage

```
write.graph(g, file = "graph.txt")
```

Arguments

g   a graph of class graphNEL

file   name of the file to be written

Value

a plain text file representing the graph. Each row corresponds to an edge represented through a pair of vertices separated by blanks

Examples

```
## Not run:
data(graph);
write.graph(g, file="graph.edges.txt");
## End(Not run)
```

Index